


Algebra and perspective

Computer Graphics (DT3025)

Martin Magnusson
November 1, 2016



Last time

- GPU programming
 - vertex shaders
 - fragment shaders
- Colour fundamentals
- Rasterised vector graphics

Vertex shader code

```
layout(location=0) in vec4 inPosition;  
out vec3 myColour;  
  
void main() {  
    gl_Position = inPosition;  
    myColour.rg = inPosition.xy;  
    myColour.b = 1.0;  
}
```

What does this vertex shader do?

- 1 Perspective correction
- 2 Apply gradient colour to vertices
- 3 All of the above

Fragment shader code

```
in vec3 myColour;  
out vec4 pixel;  
  
void main() {  
    pixel.rgb = myColour;  
    pixel.a = 0.0;  
}
```

What does this fragment shader do?

- 1 Interpolate (linearly) between vertex colours
- 2 Pin all depth coordinates to zero
- 3 Pass through colours from vertex shader


Today

- How to compute where an object ends up on the screen.
- How to check which objects occlude each other.

Reading material


- Hughes et al.:
 - 7.1–7.6.6
 - 10 (mostly 10.6 and 10.13)
 - 11.1–11.2.1
 - 13

Spaces in computer graphics



- a)** Object space (local coordinates, per model)
- b)** World space (global coordinates, complete scene)
- c)** Eye space (camera-local coordinates)
- d)** Image space (perspective)
- e)** Screen space (2D)

Transforming between spaces



- 1 object \rightarrow world space: translate and rotate from world pose
- 2 world \rightarrow eye space:
 - translate so that camera is at origin
 - rotate (around origin) so camera looks along $-z$ and y is up
- 3 eye \rightarrow image space: perspective transform (scale by $1/z$)
- 4 image \rightarrow screen space: remove z component

Vectors

3D vector $\mathbf{v} = [v_x, v_y, v_z]^T = \begin{bmatrix} v_x \\ v_y \\ v_z \end{bmatrix}$

Norm $\|\mathbf{v}\| = \sqrt{v_x^2 + v_y^2 + v_z^2}$

Addition $\mathbf{u} + \mathbf{v} = \begin{bmatrix} u_x + v_x \\ u_y + v_y \\ u_z + v_z \end{bmatrix}$

Scalar (dot) product $\mathbf{u} \cdot \mathbf{v} = u_x v_x + u_y v_y + u_z v_z = \|\mathbf{u}\| \cdot \|\mathbf{v}\| \cdot \cos \theta$

Vector (cross) product $\mathbf{u} \times \mathbf{v} = \begin{bmatrix} u_y v_z - u_z v_y \\ u_z v_x - u_x v_z \\ u_x v_y - u_y v_x \end{bmatrix} = \text{scaled normal}$

Matrices

Matrix $\mathbf{A} = \begin{bmatrix} a_{11} & a_{12} & a_{13} \\ a_{21} & a_{22} & a_{23} \\ a_{31} & a_{32} & a_{33} \end{bmatrix}$

Multiplication $\mathbf{AB} = \mathbf{C}$, with $c_{ik} = a_{i1}b_{1k} + a_{i2}b_{2k} + \cdots + a_{im}b_{mk}$

Transpose $\mathbf{A}^T = \begin{bmatrix} a_{11} & a_{21} & a_{31} \\ a_{12} & a_{22} & a_{32} \\ a_{13} & a_{23} & a_{33} \end{bmatrix}$

Inverse \mathbf{A}^{-1} , such that $\mathbf{AA}^{-1} = \mathbf{A}^{-1}\mathbf{A} = \mathbf{I} = \begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix}$

Homogenous coordinates

- Add an extra (fourth) element $[x, y, z, w]^T$
- Represents the point $[x/w, y/w, z/w]^T$
- Typically: $[x, y, z, 1]^T$

Vectors vs. points

- So, our 3D world is the slice of the 4D space where $w = 1$.
- Points in space have $w = 1$.
- Vectors have $w = 0$.
- Why?
 - Vectors have no “place” in space, but points do.
 - We can add vector + vector (= a new vector with $w = 0$), and vector + point (= a new point with $w = 1$)
 - but we can't add two points. (“This corner plus that corner” doesn't mean anything.)

Transformations

- How do we express transformations on vectors in 2D and 3D space?
- Matrix/vector multiplication is convenient.
- But we could also *add* vectors, use quaternion algebra, etc.


Scaling

$$\begin{bmatrix} s_x & 0 & 0 & 0 \\ 0 & s_y & 0 & 0 \\ 0 & 0 & s_z & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix} = \begin{bmatrix} s_x x \\ s_y y \\ s_z z \\ 1 \end{bmatrix}$$

Rotation


Rotate by angle θ :

- “around z axis,”
- equivalently: “in xy plane.”




$$\begin{bmatrix} \cos \theta & -\sin \theta & 0 & 0 \\ \sin \theta & \cos \theta & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix} = \begin{bmatrix} x \cos \theta - y \sin \theta \\ x \sin \theta + y \cos \theta \\ \textcolor{red}{z} \\ 1 \end{bmatrix}$$

Shear



2D shear along x



2D shear along y

$$\begin{bmatrix} 1 & 0 & a & 0 \\ 0 & 1 & b & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix} = \begin{bmatrix} x + az \\ y + bz \\ z \\ 1 \end{bmatrix}$$

Translation (moving)

- In Cartesian coordinates, no matrix exists that can do translation.
- We'd need to *add a vector*, not *multiply by matrix*.

$$\begin{bmatrix} t_x \\ t_y \\ t_z \end{bmatrix} + \begin{bmatrix} x \\ y \\ z \end{bmatrix} = \begin{bmatrix} x + t_x \\ y + t_y \\ z + t_z \end{bmatrix}$$

- Can we make it fit our matrix multiplication framework anyway?
- Cue: homogeneous coordinates.

$$\begin{bmatrix} 1 & 0 & 0 & t_x \\ 0 & 1 & 0 & t_y \\ 0 & 0 & 1 & t_z \\ 0 & 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix} = \begin{bmatrix} x + t_x \\ y + t_y \\ z + t_z \\ 1 \end{bmatrix}$$

Perspective

If z is “out of the screen” in eye space (so $-z$ is “into the image”), and the image plane is at $z = d$:

$$\begin{bmatrix} d & 0 & 0 & 0 \\ 0 & d & 0 & 0 \\ 0 & 0 & d & 0 \\ 0 & 0 & 1 & 0 \end{bmatrix} \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix} = \begin{bmatrix} dx \\ dy \\ dz \\ z \end{bmatrix} = \begin{bmatrix} dx/z \\ dy/z \\ d \\ 1 \end{bmatrix}$$

homogenisation

If $d = 1$:

$$\dots = \begin{bmatrix} x/z \\ y/z \\ 1 \\ 1 \end{bmatrix}$$

Combinations

- Transformations can be combined using matrix multiplication.
- *Order is important* (not commutative).

$$\underbrace{\begin{bmatrix} 1 & 0 & 0 & t_x \\ 0 & 1 & 0 & t_y \\ 0 & 0 & 1 & t_z \\ 0 & 0 & 0 & 1 \end{bmatrix}}_{\text{translate}} \underbrace{\begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & \cos \theta & -\sin \theta & 0 \\ 0 & \sin \theta & \cos \theta & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix}}_{\text{rotate around } x} \underbrace{\begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix}}_{\text{point}} = \underbrace{\begin{bmatrix} xs_x + t_x \\ ys_y \cos \theta - zs_z \sin \theta + t_y \\ ys_y \sin \theta + zs_z \cos \theta + t_z \\ 1 \end{bmatrix}}_{\text{transformed point}}$$

Rotation matrices

- A 3×3 matrix (or homogeneous 4×4) can represent all possible rotations — but all matrices are not rotations!
- Requirements:
 - square
 - $|\mathbf{R}| = +1$
 - orthogonal: $\mathbf{R}^T = \mathbf{R}^{-1}$
- In other words:
 - \mathbf{R} is normalized: the squares of the elements in any row or column sum to 1.
 - \mathbf{R} is orthogonal: the dot product of any pair of rows or any pair of columns is 0.
- The *rows* of \mathbf{R} represent the axes in the *original* space of unit vectors along the axes of the *rotated* space.
- The *columns* of \mathbf{R} represent the axes in the *rotated* space of unit vectors along the axes of the *original* space.


Orthonormalization

- What to do if rotation matrix is not orthogonal and with determinant 1?
 - If we know that it should be a rotation matrix (only), we can “massage” it into being orthonormalised again.
- 1 Normalise first row (or column).
 - 2 Cross product of first and second row, normalise, use result as third row.
 - 3 Cross product of first and third row, use result as second row.
- May not be “the correct” rotation anymore, but at least it will be a rotation.

Euler angles (fixed angles)

$$(\theta_x, \theta_y, \theta_z)$$


- Rotation orders: (x, y, z) , (z, y, x) , (x, y, x) , etc.
- Problem: gimbal lock.
- Problem: interpolation \rightarrow “detour”.



Axis and angle

$$([x, y, z], \theta)$$

- Easy to read
- No gimbal lock
- Hard to concatenate
- Non-trivial to interpolate



Quaternions

$$[s, x, y, z]$$

- (Unit-length) quaternions : $\mathbf{q} = [s, x, y, z]$
- Generalisation of complex numbers
- NB: *not* homogeneous coordinates.
- NB: *not* axis+angle.
- “Vector part” (imaginary) is $\sin(\theta/2) \cdot \text{axis}$
- “Scalar part” (real) is $s = \cos(\theta/2)$
- No gimbal lock, easy to combine, easy to interpolate

Vector products

What does the dot product between two vectors represent?

- 1 The cosine of the angle between the vectors.
- 2 A third vector, perpendicular to the two.
- 3 None of the above.

▶ [Link](#)

Matrix inverse

What is the inverse of

$$\mathbf{A} = \begin{bmatrix} 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 \end{bmatrix}$$

- 1 \mathbf{A}^{-1} does not exist
- 2 \mathbf{A}^{-1} is the identity matrix
- 3 $\mathbf{A}^{-1} = \mathbf{A}^T$

Transformations

For a point $\mathbf{x} = (10, 0)$ how to rotate it 5 degrees around $(8, 2)$?

- T_1 : translate $(8, 2)$
- T_2 : translate $(-2, 2)$
- R : rotate 5 degrees

1 $\mathbf{x}' = T_1^{-1} R T_1 \mathbf{x}$

2 $\mathbf{x}' = T_1 R T_1^{-1} \mathbf{x}$

3 $\mathbf{x}' = T_2^{-1} R T_2 \mathbf{x}$

4 $\mathbf{x}' = T_2 R T_2^{-1} \mathbf{x}$

5 $\mathbf{x}' = -T_1 R \mathbf{x}$

Rotation matrices


- Requirements for a rotation matrix:
 - square
 - $|\mathbf{R}| = +1$
 - orthogonal: $\mathbf{R}^T = \mathbf{R}^{-1}$

What happens if the determinant $|\mathbf{R}| = -1$?

- 1 \mathbf{R} is a reflection
- 2 \mathbf{R} is a rotation with a negative angle
- 3 \mathbf{R} does not exist

Handling occlusion

- How do we compute which objects should be visible?




Painter's algorithm


- 1 Sort primitives by distance to camera
 - 2 Draw most distant primitives first
 - 3 Paint nearby objects on top of old ones
- + Simple to implement
 - Expensive to sort all objects
 - Expensive to draw pixels that will be overwritten
 - Still cannot handle all scenes

z-buffer

- In addition to framebuffer, use a *z-buffer*: an array with a *z* value for every *pixel*.
- Only update pixel if new *z* is closer than the buffer's value.




A simple three dimensional scene




Z-buffer representation

z-buffer outline

- 1 Reset values in z -buffer to infinity.
- 2 Draw objects *in any order*.
- 3 When drawing, compute z value for every pixel.
- 4 Only update buffer (colour and z -buffer) if new z value is smaller than old value.



A simple three dimensional scene



Z-buffer representation

z-buffer pros and cons

Disadvantages:

- Memory usage
- Have to compute z value for all pixels
- Precision problem when many objects compete for same pixel (e.g., at edges)

Advantages:

- + Simple algorithm
- + Efficient hardware implementations
- + Works for all *nontransparent* objects / scenes

What about transparent objects, then?

- A-buffer: “the anti-aliased, area-averaged, accumulation buffer” (Carpenter 1984)
- Instead of storing single z value, build a *linked list* for each pixel.
- Fragment shader “draws” all pixels, adding to the list.
- Keep list sorted on depth.
- Post processing: traverse list, compute final colour with blending.

Carpenter's A-Buffers



Clipping

- During *primitive assembly* the vertices are *clipped* to fit a volume

$$-c \leq x \leq c$$

$$-c \leq y \leq c$$

$$-c \leq z \leq c$$

- Why clip on z axis?
 - Don't draw objects behind the camera
 - Avoid numerical problems (div by zero)
 - Avoid drawing the whole world to infinity

Transformations in OpenGL

- Compute final transformation matrix as *product* of a sequence of *primitive* transformations.
- Pass to vertex shader (as a `uniform`) variable.
- Vertex shader (typically) multiplies each vertex position with this matrix.
- In legacy OpenGL, we could use `gluPerspective` to automatically set up projection matrix.
- Since OpenGL 3.3, we need to set it up ourselves and pass it as input to the vertex shader. (E.g., the GLM or glhlib projects)
- (See Lab 1.6!)

Projection caveat

- Applying the projection matrix from before works, but...
- after flattening the scene onto the $z = 1$ plane, we lost all depth info.

$$\begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 1 & 0 \end{bmatrix} \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix} = \begin{bmatrix} x \\ y \\ z \\ z \end{bmatrix} = \begin{bmatrix} x/z \\ y/z \\ \textcolor{red}{1} \\ 1 \end{bmatrix}$$

- We will need the distance for z -buffering (to determine what is in front of what).

Making a view volume

- Instead of flattening onto $z = 1$ *plane*, make view frustum *box* instead.
- We want to map the full range of z values to $[-1, +1]$ and keep *relative distances*.
- $\text{pseudo}(z) = A + B/z$
- Choose A, B so that clipping planes are at $+1$ and -1 .
- (Near clipping plane: $z = +1$, far clipping plane: $z = -1$.)

$$\begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & A & B \\ 0 & 0 & -1 & 0 \end{bmatrix} \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix} = \begin{bmatrix} x \\ y \\ Az + B \\ z \end{bmatrix} = \begin{bmatrix} x/z \\ y/z \\ A + B/z \\ 1 \end{bmatrix}$$

Accounting for aspect ratio

- One more thing: we also need to account for non-square view volumes.
- If window aspect ratio is x/y (width/height), we'll need to scale x with y/x .

$$\begin{bmatrix} H/W & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & A & B \\ 0 & 0 & -1 & 0 \end{bmatrix}$$

(Where H/W would be $9/16$ for a normal 16:9 display.)

Transforming z to fit $[-1,+1]$

- Near clipping plane (e.g., $n = -1$) should map to $+1$.
- Far clipping plane (e.g., $f = -10$) should map to -1 .

$$1 = A - B/n$$

$$-1 = A - B/f$$

- Solve for A, B :


$$A = -\frac{f+n}{f-n}$$

$$B = -\frac{2fn}{f-n}$$

Putting it all together

- 1 Set up your model matrix (move/rotate object to where it is in the world).
- 2 Set up your view matrix (move/rotate world to camera's point of view).
- 3 Set up your projection matrix (as in the previous slide).
- 4 Combine the three to a MVP (modelViewProjection) matrix.
- 5 Pass the matrix to the vertex shader (with a uniform).


```
layout(location=0) in vec4 inPosition;  
uniform mat4 projectionMatrix;  
  
void main() {  
    gl_Position = projectionMatrix * inPosition;  
}
```



Using z-buffering in GL

- Enable z-buffer (depth test), and specify sorting (keep smaller or larger z values).

```
glEnable(GL_DEPTH_TEST); // happens after fragment shader  
glDepthFunc(GL_LESS);
```



Summary

Quaternions : compact
rotation representations

— But you'll use rotation matrices instead

Matrix multiplication, order

Homogeneous coords., and their use

Translation = high-dimensional shear

Projection

Depth buffering: Z-buffer, A-buffer

What's next

Next lecture: lighting and materials

- Mon Nov 7, 13.15–15.00
- T-141
- Hughes et al.:
 - 6.2.2–6.3, 6.5,
 - 14.9,
 - 1.13.1–2.
 - (Chapter 27 is great, but perhaps overly detailed. I recommend to look through it, but never mind solid angles and integrals for now. The chapters listed above are more to the point.)

What's next

Next next lecture: textures

- Tue Nov 8, 14.15–17.00
- T-211
- Hughes et al.:
 - 7.9–7.9.1,
 - 9.6,
 - 20.1–20.8.2.

Study questions (for lecture #3)

- 1 The irradiance (incoming light energy per area) at a surface patch is proportional to the cosine of the angle of the incoming light. Why?
- 2 How many dimensions (input and output values) does a BRDF function have?
- 3 How many dimensions does a BSDF function have?
- 4 Search for the most *Lambertian* surface you can see (if there is one).
- 5 Search for the surface with the highest *specular* exponent (if there is one).
- 6 Search for the surface with the highest *ambient* component (if there is one).

References



Loren Carpenter (1984). “The A-buffer, an Antialiased Hidden Surface Method”. In: 18.3, pp. 103–108.



John F. Hughes et al. (2013). *Computer graphics: principles and practice (3rd ed.)* Boston, MA, USA: Addison-Wesley Professional, p. 1264. ISBN: 0321399528.