

Shadows with rasterisation

Computer Graphics (DT3025)

Martin Magnusson
November 14, 2016



Doom 3, id Software



Assassin's Creed, Ubisoft

Last time

- Mapping 2D image data to 3D objects
- Interpolating over triangles (barycentric coordinates)
- Mapping applications
 - Texture maps
 - Bump maps
 - Displacement maps
 - Environment maps
- Aliasing, anti-aliasing


Today

- Tricks for casting shadows with rasterisation graphics
 - Shadow volumes
 - Shadow mapping

Reading material


- EuroGraphics'10 course notes:
 - Chapters 1, 2, 7 (and some of 3).
 - But mostly Chapter 2!

Objects/transformations/clipping




What we've done

Basic direct lighting



Where to go


Casting shadows



no ambient



with ambient

Indirect illumination (and transparency, etc)



Types of shadows

Light and shadow




Shadows give important depth cues.

Shadow components

Occluder creates shadows

Receiver is shadowed



Types of shadows

Canonical shadows

source has a finite size, a perhaps irregular shape and a definite position in space relative to the objects to be represented. Light sources of finite size cast shadows involving an umbra and penumbra. The umbra is that part of the shadow space which receives no light from the source; the penumbra, that part which receives light from some part of the source but not all of it. Thus there is a dark central area in which a smooth change from shadowed to unshadowed takes place. For an irregularly shaped light source the penumbra could be approximated by a linear variation in shade over a strip of fixed width around the periphery of any umbra. Calculating a penumbra can be expected to significantly increase the effort necessary to represent shadows. Therefore, a point light source or one infinitely far removed (specified by a direction only) will be assumed.

Shadow boundaries are determined by projecting the silhouette of one object onto another. The type of projection which may be used is dependent on the position of the light source. The nearest light source for which to calculate shadows is one that is infinitely far removed since shadow boundaries may be found by an orthographic projection. On the other hand, the calculation of shadow boundaries for a light source which has a position in the object space varies in difficulty with the location. If the source lies outside the field of view, shadow boundaries can be calculated by using the same sort of perspective projection used for image display. However, when the light source lies within the field of view, different methods must be used. Since the conventional perspective

transformation is accurate only for a limited field of view, either the space must be divided into sources radiating from the light source, in which the perspective transform from the observer, or more complicated three-dimensional geometric methods must be used.

Projective transforms provide convenience and efficiency. However, it is always possible to define shadow boundaries in the object space by using the light source position and the object silhouette to define a surface and then calculating the intersection of that surface with other objects.

CLASS ONE: SHADOW COMPUTATION USING THE SCANNING

Appel [1,2] and then Smith and Bailey [3] have shown methods for rendering shadows which calculate shadow boundaries while scanning the image. Appel generates shadow boundaries by extending his notion of quantitative invisibility. Quantitative invisibility is a count of the number of surfaces hiding a vertex (polygonal objects are assumed). Therefore, a line segment is visible only if all points on it have a quantitative invisibility of zero. Changes in quantitative invisibility along a segment are detected by Appel's hidden surface algorithm and only the visible portions are drawn. This method yields a line drawing.

Shadowed surfaces are determined during a scanning procedure which is also used to shade the line drawing. The scan is executed by generating "cutting" planes through the polygons which intersect




Figure 1: ABE defines a "cutting plane" in Appel's algorithm.

Edges of polygon 1 are projected onto polygon 2 to form shadow boundaries which are then projected onto the image plane.



fig. 1
The observer's view of four spheres.

fig. 2
The four spheres viewed from the position of the light source.



fig. 3
Shadow with reduced surface bias reveals quantization noise.

fig. 4
Increased bias and dither applied to shadow computation.



fig. 5
Low-pass filtering applied to shadow.

fig. 6
Shadow applied to the image of fig. 1

Microscopic shadows






Figure 1.2 What we define as *shadow* depends upon the scale at which we look at objects. In the real world, the definition is thus very ambiguous; in a virtual world, described by a mathematically accurate framework, precise definitions are possible and meaningful. Left: Courtesy of Prof. U. Hartmann, Nanostructure Research and Nanotechnology, Saarland University. Right: Courtesy of [flickrPrince, 2007]


- Shadows vs material properties.
- Compare with Cook–Torrance reflectance.

Types of shadows



- Dark side at the back of an object
 - OK, we can already deal with this — but remember to check if dot product is positive!



Types of shadows



- *Casting* shadows on an object
 - we'll deal with that today



Casting shadows on planar surfaces

- Draw the object a second time:
 - coloured in as $\text{RGBA} = (0, 0, 0, 0.5)$ (for example),
 - and transformed with an extra matrix that projects vertices to ground plane.
 - What does such a projection matrix look like?
 - Any limitations?




Limitations of planar shadows

- Much more difficult with curved surfaces.
- No self-shadowing.
- Shadow polygons falling outside receiver.



Shadow volumes

- Essentially, we create a new 3D shape representing the *shadow volume* of an object.
- Shadow volume of triangle: capped triangular pyramid with point light at apex.



Eisemann et al. 2010

Shadow volume test

Conceptual pseudo code (“is the point inside the shadow volume of any light and polygon?”)


```
1  bool in_shadow( point P )
2  {
3      for T in all triangles
4      {
5          for L in all lights
6          {
7              if (inside_shadow_volume( P, L, T ))
8                  return true;
9          }
10     }
11     return false;
12 }
```

Shadow volume test

- Similar to the in-triangle test from lecture 1.
- Check if each fragment is inside all the half-spaces of the shadow-volume boundaries.
- Cost of naive implementation: $O(\text{\#polygons} \times \text{\#lights})$

A slightly better shadow volume test

- Is \mathbf{p} in shadow?
- 1 Shoot ray from view point to \mathbf{p} . Set counter $c = 0$.
- 2 Increment c each time ray enters a shadow volume.
- 3 Decrement c when ray exits shadow volume.
- 4 $c = 0$: \mathbf{p} is lit.
- 5 $c \neq 0$: \mathbf{p} is in shadow.



Smarter shadow volumes

- 1 We don't need shadow volume for *each primitive*.
 - \Rightarrow Compute shadow volumes only for *silhouette edges*.




Figure 2.5 Left: Interior edge makes two quads which cancel out. Right: Finding the silhouette edges gets rid of many useless shadow volume quads.

- 2 Use the *stencil buffer* for more efficient intersection lookups.

The stencil buffer

- An additional buffer (same dimensions as framebuffer and depth buffer.)
- Typically used to mask out some pixels.

// Stencil test: write to depth/frame buffers only if test passes.

// This case: pass stencil test when stencilValue == 1.


```
glStencilFunc(GL_EQUAL, 1, 0xFF);
```

- Has operation such as increment or decrement based on result of stencil and depth tests.

```
glStencilOp(sfail,dpfail,dppass);
```

Stencil shadow volume algorithm

- 1 Draw scene with only ambient light.
- 2 Turn off frame and depth buffer.
- 3 Generate shadow volumes from triangle edges.
- 4 “Draw” *forward-facing* volumes, incrementing stencil buffer +1.
- 5 “Draw” *backward-facing* volumes, decreasing stencil buffer -1.
Where stencil $\neq 0$ we now have a shadow.
- 6 Turn on frame and depth buffer.
- 7 Draw scene again, diffuse/specular light where stencil = 0.



Note that these images show diffuse+specular also in intermediate stages.

Stencil buffer shadows

Example: Doom 3




Doom 3, id Software

z-fail

- What if the camera is also in shadow?
- Count c will be wrong.
- Solution 1: initiate c to “# shadow vols that camera is in”.
 - Need extra geometry tests on CPU.
- Solution 2: invert the depth test (known as “z-fail”)
 - Count shadow-volume fragments *behind* \mathbf{p} .
 - I.e., “shoot ray from \mathbf{p} *away from* view point.”
 - Rationale: point at infinity is always in the light.
 - “z-fail” because the counted shadow-volume fragments fail the depth test (they are not visible).
 - More robust, but slower (must update stencil buffer also for many occluded fragments)

ZP+

- Quick way to initiate the stencil buffer (so that we can apply z-pass).
- Render scene from the light onto the camera's near plane.
- Initialize stencil buffer with correct values.



The limitations of shadow volumes

- Shadow volumes give exact shadows, but...
- Extracting silhouette edges and rasterising shadow polygons becomes *infeasible* for more complex scenes.



Crysis Warhead, CryTek

Background

Shadow mapping

- So, shadow volumes are too heavy for complex scenes. What to do?
- Shadow mapping.
- Pros:
 - No explicit processing of scene geometry required!
 - Fewer rendering passes than shadow volumes for a single light source.
 - (No need to use stencil buffer.)
- Cons:
 - Aliasing artifacts may occur.
 - Needs an extra rendering pass for each additional light.

CROTTING CURVED SHADOWS ON CURVED SURFACES

Lawrence Williams
Computer Graphics Lab
New York Institute of Technology
Old Westbury, New York 11568

Abstract

Shadowing has historically been used to increase the intelligibility of scenes in television animation and video games. The first serious work was published for the determination of shadows in computer synthesized scenes. The mapping of shadows may have the shape and relative position of objects in such scenes more comprehensible if it is a technique lending realism and realism to computer animation.

To date, algorithms for the determination of shadows have been restricted to scenes composed of planar polygons. A simple algorithm is described which utilizes a hidden-surface surface computation to display shadows cast by objects modeled as smooth surface patches. The method can be applied to all environments. In fact, for each visible surface can be computed. The cost of determining the shadows projected onto each light source is roughly twice the cost of rendering the scene without shadows, plus a fixed transformation overhead which depends on the image resolution. Because shadows are not added to the scene description in the shadowing process, this representation algorithm can be used to generate shadows on curved surfaces without the need for a hidden-surface method for casting the shadows of the environment on a single ground plane.

In order to attain good results, the discrete nature of the visibility computation must be treated with care. The effects of aliasing, interpenetration, and penumbra simulation are examined. The special problems posed by self-shadowing surfaces are described.

Key words: shadows, hidden-surface algorithms, computer animation, computer graphics.

CR classifications: A.2

Introduction

The 3-buffer visible surface algorithm, first published by Catmull [1], was the first method to make possible computer generated shaded pictures of realistic surface patches. The algorithm is extremely general and quite flexible to implement but requires substantial memory.

A "frame buffer," in the current computer graphics parlance, is a memory that stores a complete digital picture. It may serve as an intermediary between the hardware that produces the picture and a frame store which eventually references a display. Some visible surface algorithms (e.g., [1]) require a frame buffer to store a complete image. In this case, the frame buffer simulates the display process in a more substantial way.

The 3-buffer is an extension of this main-memory approach to compute graphics and requires the visible surface in a scene by storing depth (z) values at each pixel. In the picture, all objects are rendered, their z values are compared at each pixel with the stored z values to determine visibility, then this determination requires only that a maximum value (depth) is stored at each pixel. In this way, the 3-buffer algorithm provides a shadowed picture to all scenes for which visible surfaces can be computed.

3-buffer visible surface computation is of particular interest because it satisfies hidden-surface projection [1]. The objects to be rendered do not have to be sorted beforehand. Individually computed surfaces are rendered. In the 3-buffer, the 3-buffer implicitly assumes pixels work in 3D and are bucket maps. The special case of the render and store the pixels occupies the frame of the buffer, limiting all computation. In the 3-buffer, the only visible surface algorithm the cost of which grows only linearly with the number of pixels in the scene is the 3-buffer.

Shadowing is the only sorting method which grows only linearly in space with the number of rendered-shaded lines to be rendered, and the 3-buffer is the only visible surface algorithm the cost of which grows only linearly with the number of pixels in the scene. In the 3-buffer, the only visible surface algorithm the cost of which grows only linearly with the number of pixels in the scene is the 3-buffer.

The 3-buffer algorithm requires the key advantages over all other existing visible surface algorithms:


1. Individually large environments;
2. Linear cost growth.

In addition, the final image computed has an appearance of realism.

Shadow map algorithm in a nutshell


First:

Render **shadow (depth) map** from the light source




During the second pass:

A fragment is in shadow if its depth is greater than the corresponding depth value in the shadow map




Shadow mapping algorithm (1/3)

- All lit parts can be “seen” by the (point) light source.
- We have efficient implementations of how to determine which parts of objects are seen from a point:
- render an image from the light source’s position!
- The *depth buffer* from this point is the *shadow map*.
- (Store it as a texture.)




Shadow mapping algorithm (2/3)

- Render the scene from the camera view as usual.
- Transform each fragment's xyz coordinates into *light space*.
- **How?**
 - We have $\mathbf{p} = (x, y, z)$ in view space; \mathbf{V} : the view matrix (brings point in world space to *view* space); \mathbf{L} : (brings point in world space into *light* space).
 - (For shadow mapping, our light source needs an orientation in addition to position.)
 - $\mathbf{LV}^{-1}\mathbf{p} = \mathbf{p}' = (x_s, y_s, z_s)$ fragment's position in light space.



Shadow mapping algorithm (3/3)



- Now, (x_s, y_s) correspond to (u, v) in the shadow map.
- z_s is the distance from the light to this fragment.
- How to determine if this fragment is occluded or visible from light source?
- Compare z_s and the depth map value d at position (x_s, y_s) .
- $z_s > d \implies$ fragment is in shadow.
- $z_s = d \implies$ fragment is lit.



Issue: Z fighting

- For points that are visible from both camera and light source:
 $\text{shadow_map}(x_s, y_s) \approx z_s$.
- Rounding errors result in random noise.
- Z-fighting, a.k.a. shadow acne.
- (Top figure shows low-resolution shadow map → aliasing.)

without depth bias

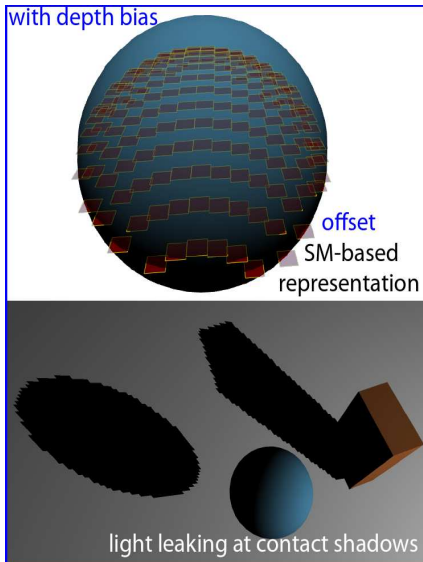


Workaround: bias

- Add bias b :


$$\text{shadow_map}(x_s, y_s) + b < z_s$$

- How much bias is good?
- Too little: shadow acne.
- Too much: *light leak*.




Depth bias: caveats

- If a face mostly aligns with the light's view, a much larger bias can be necessary.
- Usually, two parameters are available, a constant offset and an offset that depends on the alignment of the triangle with the light's view rays.
- Needs to be hand-adjusted.
- E.g., for a very short triangle, too much offset \implies shadow depth does not correlate to geometry.



Issue: Shadow map aliasing


- Shadow map has limited resolution.
- Shadow-map pixels (texels) do not correspond 1-to-1 with screen pixels.
- Workaround 1: high-res shadow maps...



Workaround 2: Percentage-closer filtering


- Countering aliasing: interpolation?
- Does not work for depth!
- Compute average of surrounding depth tests instead.

First: Render SM from light




Second: Render from view

Test pixel (reference depth)
against SM window depths



Third:

Average depth test outcome



42 blocked, 7 visible


PCF shadow = $7 / (42+7)$

0.14

Read more in Eisemann et al. 2010 if you are interested in the details about these workarounds.


Workaround 3: Fitting

- Make sure shadow map is fitted to view frustum.
- (Making sure no shadow texels are wasted.)
- Can cause *temporal aliasing* because fitting is recomputed each frame.




Workaround 4: Warping

- Perspective shadow maps.
- Transform scene (log along z axis) before projecting into shadow map.
- Leads to more uniform samples.




Without warping, fewer shadow map texels near camera.



Warp scene with $\log z$ before creating shadow map.

Workaround 5: Partitioning

- Use separate shadow maps, partitioned along view direction.



Issue: Omnidirectional light sources


- Shadow map is computed with standard rendering.
- So we are confined to a “light frustum”.

Workaround: multi-pass shadow mapping

- Need to render several (6) shadow maps.

(Planar reflections + stencil buffer)

- The stencil buffer can also be used to work around some of the limitations of planar shadows and reflections:



Left: proper stenciling. Middle: head peeks beyond surface. Right: showing that reflection is just a mirrored object.


Shadow maps or shadow volumes?

- Shadow maps are generally faster (governed by cost of rendering the image for the viewpoint)
- and can generate shadows from any rasterizable geometry (not only polygonal).
- But shadow maps have biasing issues, undersampling artifacts (jaggy edges),
- and are limited to a single frustum: so omnidirectional lights (as opposed to directional spotlights) require 6 shadow maps.
- Shadow volumes produce perfectly sharp shadows, but require 3 render passes per light, additional geometry processing, and the elongated shadow quads cause a high *fill-rate*.

Summary

- Casting shadows on other objects (with rasterisation)
- Shadow volumes
 - Gives exact hard shadows
 - Needs polygonal geometry
 - Doesn't scale well with complex geometry
- Shadow maps
 - Less sensitive to complex geometry
 - Useful also for non-polygon geometry
 - Issues with z-fighting and aliasing

Next lecture: ray tracing



- Fri Nov 18, 13.15–15.00
- T-211
- Hughes et al:
 - 15.1–15.2.4, 15.4–15.4.1, 15.4.3
 - 7.8
 - 29

References



Franklin C Crow (1977). “Shadow Algorithms for Computer Graphics”. In: *SIGGRAPH '77*, pp. 242–248.



Elmar Eisemann et al. (2010). “Shadow Algorithms for Real-time Rendering: EuroGraphics 2010 Tutorial Notes”. In: *EuroGraphics*.



Mark J. Kilgard (1999). “Advanced OpenGL Game Development”. In: Nvidia. Chap. Improving shadows and reflections via the stencil buffer, pp. 204–253. URL: <https://developer.nvidia.com/sites/default/files/akamai/gamedev/docs/stencil.pdf>.



Marc Stamminger and George Drettakis (2002). “Perspective Shadow Maps”. In: *SIGGRAPH '02*. San Antonio, Texas: ACM, pp. 557–562. ISBN: 1-58113-521-1. DOI: 10.1145/566570.566616. URL: <http://doi.acm.org/10.1145/566570.566616>.



Lance Williams (1978). “Casting curved shadows on curved surfaces”. In: *SIGGRAPH '78*, pp. 270–274.