

Real-Time Programming

Farhang Nemati


Spring 2016

Introduction

- **Multi-core Architectures**

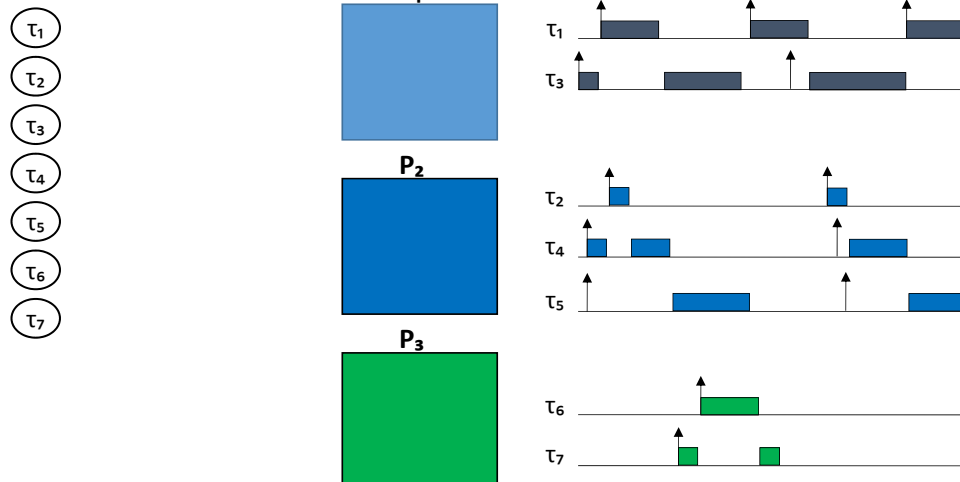
- Two or more independent processors (cores) on one chip (Single-chip multiprocessor).
- Processors may have independent (L1) and/or on-chip shared (L2) cache.
- Actual parallelism.
- Very fast inter-core communication.
- Shared Memory and BUS.
- Cores may have identical or different performance.
- Overcome thermal and power consumption problems.

Introduction



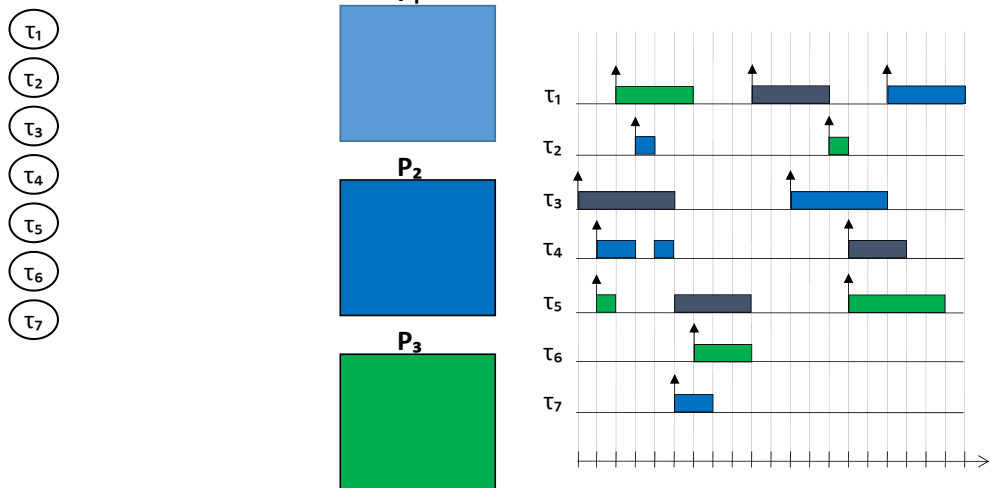
Scheduling on Multiprocessors

- Partitioned Scheduling



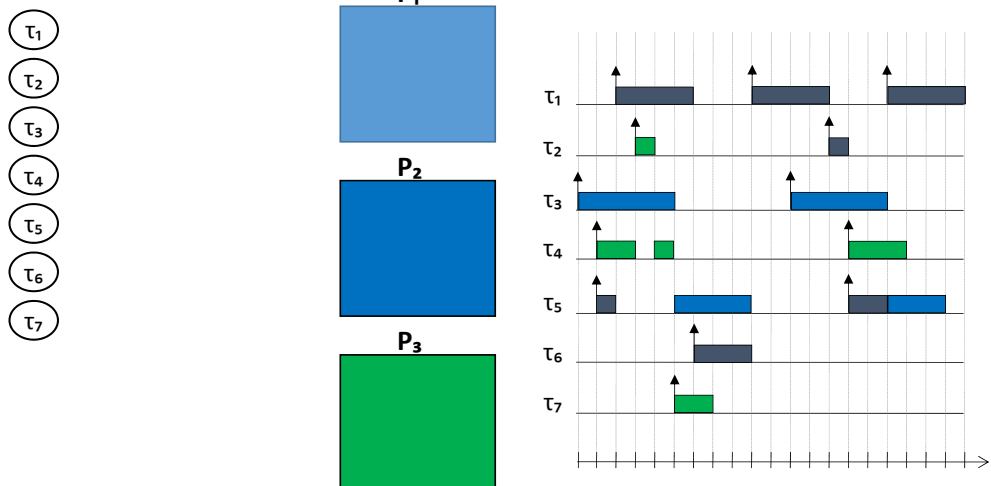
Scheduling on Multiprocessors

- Global Scheduling



Scheduling on Multiprocessors

- Hybrid Scheduling



Assign Tasks to Cores (Linux)

```
#include <pthread.h>
```

```
int pthread_setaffinity_np(pthread_t thread, size_t cpusetsize, const cpu_set_t *cpuset);
```

```
int pthread_getaffinity_np(pthread_t thread, size_t cpusetsize, cpu_set_t *cpuset);
```

Some macros for working with **cpu_set_t**:

```
void CPU_ZERO(cpu_set_t *set);
```

```
void CPU_SET(int cpu, cpu_set_t *set);
```

```
void CPU_CLR(int cpu, cpu_set_t *set);
```

```
int CPU_ISSET(int cpu, cpu_set_t *set);
```

Assign Tasks to Cores (Linux)

Set the attribute object of thread to assign the thread to specific processors at the time the thread created creating:

```
int pthread_attr_setaffinity_np(pthread_attr_t *attr, size_t cpusetsize, const cpu_set_t *cpuset);
```

Assign Tasks to Cores; Example

- See the attached code (`setAffinityExample.c`)